3. Übung IBN

Maurice Donner   Ise Glade

May 17, 2022

Aufgabe 1

Using spinlocks on a single core unit wouldn't make much sense, as a critical section cannot be accesed by two processes at the same time. Therefore there are no race conditions. Locking the only available processor on a single core (remember we are talking about spinlocks) might not be the best idea either.

Aufgabe 2

a) `wait()` unblocks a process, if `S` is larger than 0.
`signal()` increases S by 1
If those two instructions were to be exchanged, the critical section would be executed not regarding if there is a queue, since the instruction to wait has not been given yet.

b) If $S \geq 2$, Two processes would be allowed to be executed at the same time, or, if $S < 2$, just one process would be started, never giving another process the signal to start.

c) If `wait()` is missing, the critical section will just be executed without checking if there is a queue.
If `signal()` is missing, only S processes can be executed, as there is no way to increase S after a process has finished.

Aufgabe 3

a) The Assertion can be entered by executing the code of Thread A twice, then executing the code of Thread B, before executing Thread A again. This way, the condition `i == 5` is fulfilled, and the error statement can be executed. Assuming `Debug_Assert(false)` is some piece of code that shall not be executed, this is unwanted, and will only happen if the threads run in a specific order.

b) The `ManualResetEventSlim` is an object with a single boolean flag and 3 methods: `Wait()`, `Set()` and `Reset()`. When a thread calls `Wait()`, it can only continue, once the reset event is signaled. This way one can implement the dependency of one thread to another, by only giving the first `Set()` signal when a certain part of the program has been run:

```
Thread A                        | Thread B
while (true) {                  | while (true) {
    sync.Wait();                | sync.Reset()
    // Execute Second part      | // Execute Second part
}                               | sync.Set() }
```

To access the `Debug_Assert()` statement, we can use the fact, that thread B doesn't have a `Wait()` condition. After setting the reset event for the first time (`i=2`), the critical region can be entered by thread A, and then immediately stopped before the `if`-condition is reached. Thread B can then increment the counter once again, so that the condition

`i%2==1` is fulfilled. If we switch now back to thread A, the `Debug_Assert` can be executed.

c) This level can be solved by entering the critical region with thread A, then releasing the semaphore again with thread B. This shouldn't happen, because while thread A is inside of the critical region, no other process should be allowed to release the semaphore again. This is clearly not the case as thread B is allowed to release the semaphore. The implementation from the lecture looked like this:

```
wait (semaphore *S) {                  signal (semaphore *S) {
    S->count--;                            S->count++;
    if (S->count < 0) {                    if (S->count <= 0 ) {
        //add this process to S->list;         // get and remove
        block();                               // process P from S->list;
    }                                          wakeup(P);
                                           }
}                                      }
```

In C#, these routines are exchanged with `semaphore.Wait()` and `semaphore.Release()`, respectively. Semaphores are defined by `semaphore = new SemaphoreSlim(a,n)`, where `a` is the starting value of the semaphore and `n` is the maximum number of allowed processes to enter the lock at the same time. By calling `semaphore.Release(n)` the semaphore counter can be increased by `n`, allowing the same number `n` of processes to enter.

d) The goal of this level is to raise an exception, which happens when thread A tries to read a value from the queue, while it is empty. This can quickly be achieved by releasing the semaphore with thread B, then interrupting it before it can add an element to the queue.
What's abstracted is the internal operations of the queue. Instead of having to write into the buffer and then having to increment the consumer pointerm, all while taking care of the maximum buffersize, you only enqueue and dequeue items.

e) Thread 2 is responsible for adding items to the queue and also signaling all threads to wake up. We can abuse this, by letting all threads execute until they reach their `Wait()`-condition. We then add an item to the queue (thread 2) and then signal all threads to wake up. If now thread 0 finishes its routine to dequeue an item, it exits the monitor with `MonitorExit()`, and thread 1 can enter, and try to dequeue. This however will raise an exception, since the queue is empty.

Aufgabe 4

The pseudocode can be found in `pseudo_4.cpp`

Aufgabe 5

a) Channels can act like messengers between threads. If we want a thread to onlt execute a specific block of code after a condition is fulfilled we could read the channel until another thread writes into it:

Thread A	Thread B

```
var myChannel: Channel[int]
open(myChannel)

myChannel.send(1)
```

```
while true:
    if myChannel.peek == 1:
        myChannel.recv() // Blocking Call
        // Execute some code
```

This however is a spinlock and should generally be avoided since it wastes precious CPU cycles. The Nim Manual[1] also contains the procs `tryRecv` and `trySend`, which can be used to send a message to a thread directly (source: `https://nim-lang.org/docs/channels_builtin.html#Channel`)

b) The blocking call happens upon calling `myChannel.recv()` (note that the channel can also be read out without blocking by using `myChannel.peek`). To ensure that no two threads can access the channel at the same time, the `myChannel.send()` proc also blocks the channel, until the number of unprocessed items is less than a set value `maxItems`. For an unlimited queue, this value can be set to 0.

Thread A	Thread B

```
var myChannel: Channel[int]
open(myChannel)
myChannel.send()
if myChannel.tryRecv():
    // Execute some code
myChannel.send()
```

```
if myChannel.tryRecv():
    // Execute some code
myChannel.send()
```

Aufgabe 6

a) A Mutex has seperate lock and condition variables. Therefore, condition variables have no history. The condition has to be tested seperately, instead of relying on a signal. Semaphores will remember the Signals given through the Semaphore counter `S`. If a Thread broadcasts a signal, the next time another Thread calls `wait()`, it will start running immediately, regardless of when the signal was given.

b) The Mutex has to hold the lock, in order for processes not to get stuck waiting. This can happen, when a wait function is called first, then a signal runs between the time, where the wait checked for a signal, and the condition. The Thread will not see that a signal has been called, and wait forever.

---

[1]The Nim Manual: `https://nim-lang.org/docs/manual.html`

Aufgabe 7

Memory	Long term memory	Working memory
Task	Stores fundamental concepts	Stores currently needed things
Comparison	Random Access Memory (RAM)	Cache
Speed	Slow	Fast
Volatility	Long storage duration	short storage duration
Unit	Junks / Items	Junks / Items
Size	$10^9 Bytes$	$10^2 Bytes$

Memory can be moved from the working memory into the long term memory by repeating the information several times. This works especially well for memorizing vocabulary. Studying everything on one day is generally less efficient then repeating the vocabulary several times during the week.